

Clean 7-dataset benchmark: BVU vs parameter-matched MLP

